

Reactive agents

Reactive policy (strategy): simple state \rightarrow response mapping (may be *stateful*). Good for realtime.

Subsumption architecture manages conflicting behaviors. Strategy is computed from the **reward** and **transition** functions to maximize immediate *and* future rewards.

Decision process: describes knowledge of rewards & transitions. **Markov**: nondeterministic transitions; **partially observable MDP**: state is uncertain; **Q-learning**: transitions and rewards are not known. **Value** of a state:

$$V(s_i) = R(s_i, a^*) + \gamma V(T(s_i, a^*)), \quad a^* \text{ the best action, } \gamma \text{ the discount factor.}$$

The **value iteration** algorithm (initialize V_i at random, pick best action by trying all, update V_i , stop when max update is small) converges to the true value. For MDP, transition function is a probability distribution: take expected value. The **policy iteration** optimizes over the policy directly by solving for the linear equations of values for the current policy. **Q-learning**: need to learn transitions and rewards (exploration / exploitation, use $\alpha(t)$).

POMDP can be converted to deterministic **MDP** by replacing state by a **belief** function.

Deliberative agents

Deliberative architecture: explicit representation of goal states and plans. Considers only the subset of states that may be visited. **Search algorithms**: DFS (too much time spent in unpromising branches), BFS (too much memory usage), Iterative Deepening (DFS + depth limit), Branch and Bound (ignore nodes with higher cost than best found so far). With adversary: minimax, $\alpha - \beta$ pruning. **Regret**: difference in outcome between playing move i and the optimal move. In games with chance, minimize expected regret or worst-case regret. Big / chance games: evaluate horizon states with a *heuristic* and / or *Monte-Carlo sampling*. **Upper Confidence Bound** for game trees:

$$\text{minimax over } UCB_i = \frac{\text{wins}_i}{N_i} + c\sqrt{\frac{\log N}{N_i}}.$$

Factored representation: Use *situation calculus*: operators transform predicates (preconditions, add-list, delete-list). **Frame axioms** carry over predicates that were not modified. **Policy** and **reward functions** should be factored. More difficult: factoring the estimate *value* of a state (use *basis functions*).

Leat-commitment principle: delay commitment to the order in which actions are taken (modulo dependencies). Determine policy by fixpoint iteration (basis functions: solve for weights that minimize the error of the value function recurrence \equiv least squares).

Partial-order planning: discover which actions can be carried out in parallel. **GraphPlan** builds layers (time) with nodes (proposition and actions) and directed edges (in: preconditions, out: add-list). Built in polynomial time, not exhaustive but complete. **SAT**: State = vector of state variables (each with a domain), add constraints for: pre and post-conditions for operators, incompatible propositions, exclusion for operators using the same resource. Time is broken up into $2n$ discreet points (states, actions).

Multiagent systems

Multiagent modalities: centralized with shared memory, centralized or distributed with message passing, decentralized (no communication, but observe common signals). **Ontologies**: provides a shared vocabulary for heterogeneous agents. **Self-interested** agents: always act to maximize their own interest. We need to create incentives. **Centralized planning**: can explicitly detect conflicts and synergy. E.g. **blackboard system** (shared memory), **publish-subscribe system** (notify on resource usage). Using reactive agents: optimize sum of rewards or each agent optimizes its own (find a NE using policy iteration, but probably not the best one). Using deliberative agents: find conflicting resource usages and similar goals. **Partial-global planning**: goal tree in which each agent inserts its partial plans (expressed with predicates). Can then reorder actions, exchange tasks. **Contract net**: cooperation protocol. Managers divide tasks, contractors place bids, contract is made for lowest bid (no negotiation). Tasks can be subcontracted. Market-based variation: managers increase prices until they obtain a solution.

Distributed multiagent systems

Distributed contract nets: make contracts asynchronously and contact agents directly. Agents try to sub-contract tasks with high *marginal cost*. Problems: incremental bidding can't work, impossible to resolve conflict, bidders must speculate on future tasks. **Distributed SAT**: variables (tasks), domains (resources that can carry out the task), constraints (timing, preconditions and resources), relations (inequality). **Sync backtracking**: each agent extends the partial solution of the previous, allows use of heuristics (async: all in parallel, but exponential number of messages needed). **Distributed DP**: organize agents in a rooted tree + backedges. Send util (constraint) messages up, parent decides the best value locally, get value messages down. Variables are collapsed, only local node knows which value to set. **Distributed local search**: start with random assignment, make local change which most reduces the number of conflicts. Neighborhood = variables connected through constraints. Changes can be async as long as there is only 1 change per neighborhood. **Breakout algorithm**: extension of min-conflict, assigning *dynamic priority* to each constraint. Pick change that reduces most the sum of priorities. All remaining conflicts have priority increased when reaching local minimum, then restart. Termination: when time count is larger than the distance to the furthest agent.

Game theory

Games: zero / general sum, 2 / N players. **Strategies**: strictly / (very) weakly dominant, pure, mixed, *minimax* (for zero-sum: *value* of the game = expected gain v_A = expected loss v_B). Computing minimax for B: find set of p_j^B such that expected gain v of A is minimized and $\forall a_i^A, \sum_{a_j^B} p_j^B R_A(a_i^A, a_j^B) \leq v$. **Utility function**: can encode risk-profile. **Support**: set of

actions with $p_a > 0$. **Nash equilibria**: no player gains from changing strategy, all other things being equal. Theorem: every game has at least a set of (mixed) NE. Compute them by removing dominated actions and going through possible subsets of actions (*supports*). For N players, NE exists but not necessarily unique or minimax. **Uncertain utilities**: *ex-ante* (no knowledge at all), *ex-interim* (own agent type is known), *ex-post* (strongest, knowledge of all types). **Bayes-NE**: NE over ex-ante expected utilities. **Ex-post NE**: strategies give highest utility no matter the value of the unknown information (not always possible).

Agent cooperation, negotiation

Correlated equilibrium: act depending on an external factor (e.g. coin flip). **Mediator**: vehicle to enforce a contract. Plays depending on the number of agents which chose the mediator. **Negotiation**: make, accept or refuse offers to agree on a joint strategy. **Strategic** or **axiomatic**. **Alternating offers**: broken, first offer has more power. Can introduce time constraints, discount factor. **Nash bargaining solution**: maximize the product of utility gains (from *conflict deal*), concession is made by deal with lowest product. Mixed strategies: can bargain about probabilities. **Monotonic concessions**: reaches NBS. Agent with lowest risk tolerance $\frac{u_i(D_i) - u_i(D_j)}{u_i(D_i) - u_i(D_c)}$ must make an offer. **Stackelberg games**: decisions in sequence between a leader and a follower.

Mechanism design

Social choice: constructing a joint preference order reflecting individual, private orderings. Implementations: dominant equilibrium, Bayes-Nash equilibrium. **Mechanisms**: social choice function + payment rule. Goal: *incentive compatibility* (best strategy for agents leads to optimizing the social choice function). **Revelation principle**: for any mechanism, there is a truthful mechanism with the same outcomes and payments (proof by construction). **VCG tax**: maximizes the sum of declared valuations. $\text{tax}(A_i) = \sum_{A_j \neq i} v_j(d_{-i}) - v_j(d_{\text{all}})$. It is truthful, rational and incentive compatible. Generalization: *Groves mechanism* \approx VCG with offset (e.g. refund part of the tax). Problems: collusions, not Pareto-efficient (tax must be wasted). **Roberts' theorem**: affine maximizer social choice functions are the only ones that can be implemented for unrestricted preference profiles with incentive compatibility. **Median rule**: IC for single-peaked preferences.

Auctions

Auctions: social choice, goals are *optimal allocation*, *individual rationality*. **Values**: *private*, *common* (entirely dependent on other's value), *correlated*. **Auction protocols**: open-cry (*Dutch*, *English*); sealed-bids (*Discriminatory*, *Vickrey*). Problems: *manipulability* (collusion, demand reduction lie), *risk-averse* or non-truthful / irrational bidding, *timing*, *side-constraints* (e.g. procurement), authenticity,

privacy. **Revenue equivalence theorem**: all 4 settings yield equivalent revenue, but not optimal allocation (not true for *correlated values*).

Multi-unit Vickrey: each agent pays the price of the bid it displaced from the set of winning bids. **Double auctions**: sort buy & sell bids in opposing orders, price is last match. M^{th} highest price is incentive-compatible for sellers, $(M+1)^{st}$ for the buyers. **McAfee**: price is average of last (sell, buy) bid, but block that one \implies loss of 1 trade but gain incentive compatibility. May bid with *price-quantity graphs* (then, sum up curves and match demand). **Bargaining**: for 1 buyer and 1 seller, NE at price $0.5 \times (b_1 + b_2)$. **Independent English Auctions** (k run in parallel): bidding strategies include *straightforward* (ignore complementarities, bid only for the best combination); *sunk-aware* (discount perceived price of won items by some factor since they cannot be returned); *price prediction*. **GVA auction** (\approx VCG tax): pay for the difference between the best allocations without and with your bid. **GSP auction** (internet ads): i^{th} highest bidder pays $(i+1)^{st}$ price to get i^{th} slot (not incentive compatible, but stable prices and high revenue).

Coalitions

Coalitions: when *side-contracts* (utility redistribution) are allowed, coalition utility *geq* others \implies stable. **Core** of a game: set of payoff distributions for which the grand coalition is stable (often empty). Core is known to be nonempty for *Superadditive* and *Convex* games ($v(S \cup T) \geq v(S) + v(T) - v(S \cap T)$). **Shapley value**: unique vector (if core is nonempty, SV payoffs is in it). $SV(a_i)$ = average value of added payoff when added that agent to a sub-coalition (over all orders). Agents not in any carrier coalitions has $SV(a_i) = 0$. **Weighted graph games**: agents are nodes, self-edges are payoffs, edges are payoffs for 2-coalitions (not possible for all games). Value of a coalition is the sum of edge weights in the subgraph.

Voting protocols

Manipulability: non-truthful voting; removing a candidate can reverse the order; vote organizer can determine the winner by changing the order in which alternatives are presented. **Condorcet winner**: alternative that beats (or ties) all others in a pairwise majority vote (doesn't always exist; in a majority graph, it is the node with only outgoing edges). **Plurality voting**: vote for your single preferred alternative (variant: carry out $n-1$ rounds, eliminate the least preferred at each round). **Borda count**: give $(n-1) \dots 0$ points to alternatives. **Slater ranking**: vote between every pair of alternatives, pick the smallest transformation to obtain a majority graph. **Gibbard-Satterthwaite theorem**: any deterministic voting protocol (≥ 3 alternatives) has one of these properties: dictatorial, non-truthful, or \exists a candidate that cannot win.

Credits

Most content taken from the lecture notes of Boi Falting's Intelligent Agents class at EPFL, 2015.

Rendered December 13, 2015. Written by Merlin Nimier-David. This work is licensed under the Creative Commons Attribution-ShareAlike 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-sa/3.0/> or send a letter to Creative Commons, 444 Castro Street, Suite 900, Mountain View, California, 94041, USA.